

1. Introduction

The purpose of this project is to implement a Retrieval-Augmented Generation (RAG) system that answers user queries based on the contents of a PDF document. The system uses LangChain, Google Gemini API, and Streamlit to provide an interactive chatbot experience.

Key capabilities:

- Answers user queries from text and tables inside PDFs.
- Uses embeddings + retrieval for context-aware answers.
- Provides a web-based interface for users via Streamlit.

2. Problem Statement

Traditional chatbots rely solely on their pretrained knowledge and cannot answer questions about specific documents.

This project solves that problem by:

- Loading a PDF into a retrievable format.
- Creating embeddings of document chunks.
- Using a retriever + LLM pipeline to answer user queries.

3. System Architecture

3.1. Components

1) Frontend (UI):

- Implemented with Streamlit.
- Provides a chat-style interface for user queries and bot responses.

2) Backend (RAG pipeline):

- **LangChain** for chaining steps.
- **Google Gemini (Generative AI API)** as the LLM.
- **Vector index** built from document embeddings for retrieval.

3) **Data Source:**

- PDF file(s) uploaded into the system.
- Example: [AAPL Q3 2022 Report](#).

3.2. **Workflow**

Step 1: Load Document

- PDF is read and split into text chunks.

Step 2: Embed & Index

- Embeddings are generated for each chunk.
- Stored in a vector index for similarity search.

Step 3: User Query

- User inputs a question in Streamlit.

Step 4: Retrieval

- The retriever finds the most relevant document chunks.

Step 5: LLM Response

- Retrieved context + user question are sent to Gemini LLM.
- The LLM generates a final answer.

Step 6: Display

- Response shown in Streamlit chat UI.

4. Implementation

4.1. Directory Structure

RAG-DocsChatBot/

```
|— DocsChat/
|   |— DocsChat.py      # Core RAG pipeline
|
|— sample_docs/
|   |— assessment.pdf   # dataset PDF
|
|— main.py              # Streamlit frontend
|— requirements.txt     # Dependencies
|— .gitignore
|— README.md
|— docs/
|   |— Report.pdf      # This documentation
```

4.2 Core Code Functions

- **initial_load():** Loads PDF, splits into chunks, creates embeddings + retriever.
- **query(question):** Handles user queries by passing them through the retrieval + LLM pipeline.
- **main.py:** Implements Streamlit UI, calls query() to display results interactively.

5. Installation & Usage

5.1. Clone Repository

```
git clone https://github.com/siikum/RAG-DOC-CHAT.git
```

```
cd RAG-DOC-CHAT
```

5.2. Create Virtual Environment

```
python -m venv venv
```

```
source venv/Scripts/activate # Windows
```

```
source venv/bin/activate # Linux/Mac
```

5.3 Install Dependencies

```
pip install -r requirements.txt
```

5.4 Add API Key

- Get a free **Google Gemini API Key** from [AI Studio](#).
- Save it in .env file:
- GOOGLE_API_KEY=your_api_key_here

5.5 Run the Application

```
streamlit run main.py
```

6. Limitations

- Works only with text-based PDFs (not scanned images without OCR).
- Limited to the quality of the retriever and embeddings.
- Currently processes one PDF at a time.
- Figures/images are not processed.

7. Future Improvements

- Support multiple PDFs at once.
- Improve table extraction (current table parsing depends on PDF text encoding).

8. Conclusion

This project demonstrates how RAG systems can effectively answer domain-specific questions using text and table information from PDF documents. It combines LangChain, Google Gemini, and Streamlit to create a practical, interactive chatbot for document-based QA.